



PDF Download
3696630.3728535.pdf
03 April 2026
Total Citations: 2
Total Downloads: 1324

Latest updates: <https://dl.acm.org/doi/10.1145/3696630.3728535>

RESEARCH-ARTICLE

RAG or Fine-tuning? A Comparative Study on LCMs-based Code Completion in Industry

CHAOZHENG WANG, Chinese University of Hong Kong, Hong Kong, Hong Kong

ZEZHOU YANG, Tencent, Shenzhen, Guangdong, China

SHUZHENG GAO, Chinese University of Hong Kong, Hong Kong, Hong Kong

CUIYUN GAO, Chinese University of Hong Kong, Hong Kong, Hong Kong

TING PENG, Tencent, Shenzhen, Guangdong, China

HAILIANG HUANG, Tencent, Shenzhen, Guangdong, China

[View all](#)

Open Access Support provided by:

[Chinese University of Hong Kong](#)

[Tencent](#)

Published: 28 July 2025

[Citation in BibTeX format](#)

FSE Companion '25: 33rd ACM
International Conference on the
Foundations of Software Engineering
June 23 - 28, 2025
Trondheim, Norway

Conference Sponsors:
[SIGSOFT](#)

RAG or Fine-tuning? A Comparative Study on LCMs-based Code Completion in Industry

Chaozheng Wang

Chinese University of Hong Kong
Hong Kong, China
adf111178@gmail.com

Zezhou Yang

Tencent
Guangzhou, China
zezhouyang@tencent.com

Shuzheng Gao

Chinese University of Hong Kong
Hong Kong, China
szgao23@cse.cuhk.edu.hk

Cuiyun Gao*

Chinese University of Hong Kong
Hong Kong, China
cuiyungao@outlook.com

Ting Peng

Tencent
Guangzhou, China
sakurapeng@tencent.com

Hailiang Huang

Tencent
Guangzhou, China
eraserhuang@tencent.com

Yuetang Deng

Tencent
Guangzhou, China
yuetangdeng@tencent.com

Michael Lyu

Chinese University of Hong Kong
Hong Kong, China
lyu@cse.cuhk.edu.hk

Abstract

Code completion, a crucial practice in industrial settings, helps developers improve programming efficiency by automatically suggesting code snippets during development. With the emergence of Large Code Models (LCMs), this field has witnessed significant advancements. Due to the natural differences between open-source and industrial codebases, such as coding patterns and unique internal dependencies, it is a common practice for developers to conduct domain adaptation when adopting LCMs in industry. There exist multiple adaptation approaches, among which retrieval-augmented generation (RAG) and fine-tuning are the two most popular paradigms. However, no prior research has explored the trade-off of the two approaches in industrial scenarios.

To mitigate the gap, we comprehensively compare the two paradigms including Retrieval-Augmented Generation (RAG) and Fine-tuning (FT), for industrial code completion in this paper. In collaboration with Tencent's WXX department, we collect over 160,000 internal C++ files as our codebase. We then compare the two types of adaptation approaches from three dimensions that are concerned by industrial practitioners, including effectiveness, efficiency, and parameter sensitivity, using six LCMs. Our findings reveal that RAG, when implemented with appropriate embedding models that map code snippets into dense vector representations, can achieve higher accuracy than fine-tuning alone. Specifically, BM25 presents superior retrieval effectiveness and efficiency among studied RAG methods. Moreover, RAG and fine-tuning are orthogonal and their combination leads to further improvement. We also observe that

RAG demonstrates better scalability than FT, showing more sustained performance gains with larger scales of codebase. Our findings provide actionable guidance for choosing and implementing appropriate methods to adopt LCMs based on specific industrial scenarios and requirements.

CCS Concepts

• **Software and its engineering** → **Software development techniques**;

Keywords

Large Code Models, Retrieval Augmented Generation, Fine-Tuning

ACM Reference Format:

Chaozheng Wang, Zezhou Yang, Shuzheng Gao, Cuiyun Gao, Ting Peng, Hailiang Huang, Yuetang Deng, and Michael Lyu. 2025. RAG or Fine-tuning? A Comparative Study on LCMs-based Code Completion in Industry. In *33rd ACM International Conference on the Foundations of Software Engineering (FSE Companion '25)*, June 23–28, 2025, Trondheim, Norway. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3696630.3728535>

1 Introduction

Code completion stands as a fundamental and critical task in code intelligence, aiming to predict subsequent code tokens or statements that developers intend to write. By anticipating developers' coding intentions, code completion remarkably enhances programming productivity and has become an indispensable feature in modern Integrated Development Environments (IDEs) [5]. The advent of Large Language Models (LLMs) [2] has led to the emergence of Large Code Models (LCMs) [16, 18], which bring substantial improvements to code completion performance [18]. This advancement has catalyzed the development of multiple commercial tools [5, 6], offering personalized code completion services.

Despite LCMs' impressive performance in general development scenarios such as utility scripts and application development, their effectiveness substantially deteriorates when applied to industrial proprietary codebases [1]. This performance gap arises from the

* Cuiyun Gao is the corresponding author.



inherent complexity and massive scale of industrial code, which often undergoes extensive business-specific implementation. For instance, our industrial partner Tencent WXG operates WeChat, China's largest instant messaging platform with over one billion monthly active users. Beyond basic communication services, the platform supports a rich ecosystem of features including social networking (Moments) and short-form videos. To accommodate this massive user base and diverse functionality demands, WXG has developed an extensive and sophisticated backend infrastructure with unique architectural patterns and optimizations. These specialized implementations significantly deviate from conventional open-source solutions, creating a substantial distributional shift between industrial code patterns and the open-source training data used by LCMs. Moreover, the sensitive nature of proprietary code and associated privacy concerns preclude the direct utilization of commercial code completion solutions, even those with superior capabilities, necessitating alternative approaches for industrial code completion.

Given the distribution gap between industrial and open-source codebase, two predominant approaches have emerged to adapt LCMs to new domains: Retrieval-Augmented Generation (RAG) [38, 65] and fine-tuning (FT) [27, 63]. Specifically, given a novel source code repository previously unseen by the model, RAG enhances generation by retrieving relative code snippets from the repository as prior knowledge. This retrieval-based augmentation manifests in two forms: (1) prompt-based augmentation [72], where retrieved similar code segments are prepended to the input query, providing contextual examples; and (2) logits-based augmentation [24, 56], where the distribution of similar usage patterns and outputs serves as a prior to calibrate the model's prediction distribution. In contrast, FT directly optimizes the model's parameters on the target codebase, aiming at aligning the model's internal distribution with the domain-specific patterns present in industrial repositories.

While both RAG and FT approaches have demonstrated effectiveness across diverse domains, such as document QA and code intelligence [27, 38, 64, 65], these studies predominantly focus on public datasets containing general knowledge. Industrial codebases present unique challenges with their complex dependencies, company-specific patterns, and proprietary frameworks, making the direct application of existing findings uncertain. The efficacy of these approaches in such specialized, real-world environments remains largely unexplored. Existing RAG-based code completion research [35, 66] has primarily focused on Python and Java codebases, where implementations tend to be more concise and less intricate than C++ systems which are widely used in Tencent WXG [49, 70]. Moreover, these studies have been conducted on relatively small-scale open-source repositories, making their findings potentially less applicable to industrial environments. The massive scale of industrial codebases significantly increases retrieval complexity, as the search space expands exponentially and the identification of relevant code snippets becomes more challenging. For instance, logits-based retrieval methods [56, 74] provide feasible and effective solutions for small codebases. However, they become prohibitively expensive when scaling to industrial repositories with hundreds of thousands of files, as they require storing and searching through massive logit matrices for each code segment. These distinctive

characteristics of industrial codebases necessitate a thorough investigation of existing adaptation approaches tailored to these specific challenges.

In this paper, we conduct a comprehensive empirical study comparing RAG and FT approaches for adapting LLMs to industrial code completion tasks. We use the backend codebase with more than 160,000 files of Tencent WXG as our experimental testbed, which contains various business functionalities and represents the complexity typical of industrial software systems. To systematically investigate adaptation approaches, we construct a subset of 120,000 C++ files as both the retrieval corpus for RAG and the training data for FT experiments. For the RAG approaches, we explore two main categories: (1) similarity-based retrieval [74], which identifies syntactically or semantically similar code snippets, and (2) dependency-based retrieval [36], which leverages project-specific dependency relationships and usage patterns. Our evaluation comprehensively compares these RAG variants against FT approaches across multiple dimensions, including effectiveness, efficiency, and parameter sensitivity.

Through extensive experiments, our study reveals several key findings in enhancing LCM-based industrial code completion. RAG and FT substantially improve LCMs' performance, with RAG achieving a higher performance ceiling. Notably, among the evaluated RAG approaches, BM25 demonstrates superior effectiveness while maintaining computational efficiency comparing with neural embedding models. The combination of RAG and FT yields further performance gains, suggesting their complementary nature. Regarding computational costs, FT primarily requires resources during the preparation stage for model training, while RAG's overhead mainly occurs during inference as it needs to retrieve and process relevant code snippets in runtime. Furthermore, we observe that RAG demonstrates better scalability than FT, showing more sustained performance gains with larger scales of codebase.

Our key contributions are summarized as follows:

- We present the first comprehensive study of LCM adaptation approaches on a large-scale proprietary industrial codebase. Our evaluation reveals the performance limitations of current state-of-the-art LCMs when directly applied to industrial code completion tasks, highlighting the necessity for more effective domain adaptation.
- We conduct extensive experiments comparing RAG and FT approaches across multiple dimensions, including effectiveness, efficiency, and data scaling. This systematic evaluation provides a detailed understanding of the trade-offs between different adaptation approaches in industrial settings.
- We derive practical insights and guidelines for developers seeking to adapt LCMs for code completion in proprietary codebases, offering empirically validated recommendations for implementing appropriate adaptation approaches.

2 Research Methodology

2.1 Data Preparation

2.1.1 Data Collection and Filtering. We collect a comprehensive snapshot of our backend C++ codebase from our department's repository. To ensure data quality and minimize potential training artifacts, we apply the following filtering criteria:

Table 1: Statistics of our dataset, where the number of tokens are calculated by the tokenizer of DeepSee-Coder.

Split	# Files	# Tokens	# LOC	# Instances	Avg. Input Token	Avg. Target Token
Train	120,000	401M	25.5M	-	-	-
Validation	23,746	69M	3.8M	-	-	-
Test	20,000	63M	3.5M	50,000	246.75	17.01

- Remove duplicate files and auto-generated framework code (i.e., compilation files) to prevent over-fitting of templated patterns
- Filter out files containing extensive Chinese comments to maintain consistency in the training corpus
- Eliminate files with unusually long C++ definitions (# define), as these rare patterns could potentially lead to catastrophic repetition in model outputs.

After applying these filtering rules, we obtained a final dataset of 163,746 valid C++ source files. We further processed these files using ClangD [4] to standardize code formatting, including consistent indentation and line break rules.

2.1.2 Dataset Splitting. To create a robust evaluation benchmark, we select 20,000 C++ files from our filtered dataset for generating input-output pairs. We employ a sliding window approach, where each input consists of 20 consecutive lines of code, and the target is the subsequent line. To ensure meaningful evaluation metrics, we apply additional filtering criteria to the target lines: (1) exclude instances where the target line contains only a single symbol, and (2) remove cases where the target line consists solely of comments.

After this process, we randomly sample 50,000 high-quality test instances, which form our evaluation benchmark for assessing model performance. The resulting test set covers various code completion scenarios while maintaining practical relevance to real-world development tasks.

For the remaining 143,746 source code files, we randomly select 120,000 for building retrieval database and fine-tuning LCMs, and 23,746 for validation. The statistics are shown in Table 1.

2.2 Retrieval Augmented Generation

2.2.1 Retrieval Database Construction. Direct utilization of our 120,000 C++ files for retrieval presents significant challenges. First, using entire C++ files as retrieval units is impractical due to their substantial length, with our codebase averaging over 3,300 tokens per file. Second, files often contain extraneous information that could potentially interfere with the model’s inference process.

To address these limitations, we decompose the files into fine-grained retrieval units. Specifically, we employ the Tree-Sitter [58] to decompose each file into more granular units, specifically extracting individual functions and classes. We treat C++ files that either lack function definitions or contain only function declarations as single retrieval units. This decomposition strategy results in a final retrieval database of 914,667 items, providing a more focused and manageable collection of code segments. Specifically, we explore two kinds of RAG methods for code completion including similarity-based and dependency-based retrieval.

2.2.2 Similarity-Based Retrieval. Given our retrieval database D containing 914,667 items, we implement a similarity-based retrieval

mechanism. Let $\mathcal{M}$ denote the embedding model. For each item $d_i \in D$, we compute its embedding vector $\mathbf{v}_i = \mathcal{M}(d_i)$. These embeddings form our vector database $V = \mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$. During inference, for a given query q , we compute its embedding $\mathbf{q} = \mathcal{M}(q)$.

The similarity score between the query and each item in the database is calculated using cosine similarity: $\text{sim}(q, d_i) = \frac{\mathbf{q} \cdot \mathbf{v}_i}{\|\mathbf{q}\| \|\mathbf{v}_i\|}$.

We select the top- K items with highest similarity scores: $d_{(1)}, d_{(2)}, \dots, d_{(K)}$, where (i) denotes the index of the i -th most similar item. Following the work [12], these K items are concatenated in ascending order of similarity and prepended to the query as examples. The augmented prompt is then fed to the model for generation:

$$\text{prompt} = d_{(K)} \oplus d_{(K-1)} \oplus \dots \oplus d_{(1)} \oplus q \quad (1)$$

where $\oplus$ denotes concatenation.

2.2.3 Dependency-Based Retrieval. Besides similarity-based retrieval, we also explore a dependency-based retrieval method that leverages function call relationships to provide relevant context for code completion. Given a code completion query q , we first extract a sequence of function calls $C = [c_1, c_2, \dots, c_m]$ using Tree-sitter [58], a robust parsing framework. For each function call c_i , we locate its corresponding function definition d_i in the retrieval codebase, resulting in a sequence of function definitions $D = [d_1, d_2, \dots, d_m]$.

The retrieved function definitions are then concatenated and prepended to the original query q to form the augmented prompt:

$$\text{prompt} = d_m \oplus d_{m-1} \oplus \dots \oplus d_1 \oplus q \quad (2)$$

This approach ensures that the model has access to the implementation details of the functions being called in the query, providing context for more accurate code completion. By explicitly including function definitions, the model can better understand the expected behavior and usage patterns of the called functions.

2.3 Fine-Tuning

For the fine-tuning approach, we conduct training at the file level using source code. Following previous work [10, 63, 69], we first tokenize and concatenate the code text into a continuous sequence of tokens $t_1, t_2, \dots, t_n$. Given a context length L , we segment the token sequence into non-overlapping blocks:

$$(t_1, \dots, t_L), (t_{L+1}, \dots, t_{2L}), \dots, (t_{n-L+1}, \dots, t_n) \quad (3)$$

The model is then trained in an autoregressive manner to predict the next token given all previous tokens in the sequence:

$$\mathcal{L} = - \sum_{j=1}^{L-1} \log P(t_{j+1} | t_1, \dots, t_j) \quad (4)$$

where $\mathcal{L}$ denotes the loss function and j means the token index.

3 Experimental Setup

3.1 Selected LCMs

In this paper, we select two kinds of popular and state-of-the-art LCMs with their versions in different sizes. In specific, our selected LCMs are:

- **DeepSeek-Coder** (DSC) [16] is trained from 2T tokens from scratch, achieving state-of-the-art performance in a variety of code intelligence tasks. Specifically, we choose DeepSeek-Coder Base in sizes of 1.3B and 6.7B in this paper.
- **Qwen2.5-Coder** (QC) [18] is a family of LCMs based on models of the Qwen2.5 series [73]. Qwen2.5 undergoes four-stage training with a total number of 5.5T tokens, presenting state-of-the-art performance in both code generation and reasoning. We choose 0.5B, 1.5B, 3B, and 7B versions of Qwen2.5-Coder.

3.2 Evaluation Metrics

Following previous work [40, 45, 59], we select accuracy, edit similarity, and BLEU score to evaluate the effectiveness of different methods.

3.2.1 Exact Match. Exact match (EM) measures the percentage of predictions that exactly match the ground truth. For code completion, this is a strict metric that requires the model to generate exactly the same code line as the target.

3.2.2 Edit Similarity. Edit similarity (ES) measures how close the predicted code is to the ground truth based on the Levenshtein distance. It provides a more nuanced evaluation than accuracy by considering partial matches.

3.2.3 BLEU Score. BLEU (Bilingual Evaluation Understudy) score [46] evaluates the quality of the generated code by comparing the n-gram overlap between the prediction and the ground truth.

3.3 Embedding Models

Following previous work [12], we utilize the following retrieval methods for similarity-based RAG:

- (1) **Random:** For each code completion query, we randomly select K items from the retrieval database.
- (2) **BM25** [41, 53]: A classical information retrieval algorithm that ranks documents based on term frequency and inverse document frequency, treating code snippets as text documents, and computing relevance scores using keywords. BM25 has been widely used for code intelligence tasks [9, 11, 75].
- (3) **CodeBERT** [8]: A BERT-based pre-trained model specifically designed for programming language understanding, which we use to compute semantic embeddings of code snippets for similarity-based retrieval.
- (4) **UniXcoder** [15]: A unified cross-modal pre-trained model for programming language, which supports both natural language and programming language understanding tasks, serving as our embedding generator for code retrieval.
- (5) **CodeT5** [67]: An encoder-decoder model pre-trained on programming languages, which we utilize its encoder to

Table 2: Hyper-parameter settings.

Hyperparameter	Value	Hyperparameter	Value
Optimizer	AdamW [19]	Warm-up steps	100
Learning rate	5e-6	Training batch size	32
LR scheduler	Cosine Scheduler [37]	Validation batch size	32
Sequence Len.	4,096	Adam epsilon	1e-8
Max. gradient norm	1.0	Precision	BF16
Max Gen. Tokens	512	Top-P	0.95

generate embeddings for measuring code similarity in our retrieval system.

- (6) **CoCoSoDa** [55]: A contrastive learning-based code search approach that employs a soft data augmentation and momentum mechanism. It dynamically masks or replaces tokens with their types for input sequences and uses multimodal contrastive learning to align code-query pairs.

3.4 Research Questions

In this paper, we aim to answer four research questions.

RQ1: How effective are RAG and fine-tuning approaches in industrial code completion? We evaluate the performance of directly applying LCMs to our code completion scenario and assess the improvements achieved through both RAG and fine-tuning approaches independently.

RQ2: Can fine-tuning be effectively combined with RAG? We investigate whether combining fine-tuning with RAG can yield better performance than either approach alone, and analyze how fine-tuned models interact with the retrieval mechanism compared to base models.

RQ3: How efficient are RAG and fine-tuning? We analyze the computational efficiency of each approach, measuring training and inference times, memory usage, and other resource requirements to understand their practical implications in industrial settings.

RQ4: How do key parameters affect the performance of RAG and fine-tuning approaches? We conduct parameter sensitivity analysis from two aspects: (1) the impact of retrieved context size (Top-k) in RAG methods on model performance, and (2) the relationship between codebase scale and model effectiveness.

3.5 Implementation Details

In the RAG framework, we implement two distinct retrieval strategies. For similarity-based retrieval, we leverage the BM25S library to construct and maintain a lexical-based retrieval database, while utilizing the Qdrant [50] service to establish a vector-based retrieval system. For dependency-based retrieval, we first employ Tantivy [57] to construct a mapping mechanism that links function names to their corresponding definitions. Subsequently, we utilize Tree-sitter [58] to analyze the code structure and extract function call dependencies, enabling us to retrieve the relevant function definitions based on the identified function names.

All the experiments are run on a server with Intel 8374C CPU and 8*A100 GPUs with 80GB of graphic memory. The hyper-parameter setting of the tuning procedure is listed in Table 2 following previous work. We enable the gradient checkpointing technique [3] for LCMs to reduce the GPU memory consumption.

For fast inference, we utilize vLLM [26] based on PagedAttention to improve efficiency. To eliminate the influence of random

sampling, we utilize greedy decoding strategy during inference. In addition, we employ the Flash-Attention technique [7] for long-context optimization.

4 Experiment Results

4.1 RQ1: Effectiveness of RAG and Fine-Tuning

In this section, we present the performance of LCMs in our industrial code completion task in Table 3. From the results, we can reach the following observations.

LCMs demonstrate suboptimal performance on our industrial code completion task. Among the six evaluated models, even the best-performing Qwen2.5-Coder 7B achieves only approximately 28% EM and 44% BLEU score. These relatively low metrics indicate that these models struggle with our complex, domain-specific codebase that differs substantially from their training distribution. Furthermore, we observe a clear correlation between model size and completion performance. As the parameter count increases, models demonstrate better capability in capturing code patterns, leading to better task performance. This suggests that while the domain gap poses a notable challenge, larger models' enhanced capacity allows them better to generalize their learned programming patterns to unfamiliar domains.

Both RAG and fine-tuning demonstrate effectiveness in enhancing model performance on industrial code completion tasks. For RAG methods, our experiments show that similarity-based retrieval (RAG_{sim}) yields substantial improvements in completion accuracy. Specifically, when using different embedding methods for retrieval, the model's EM increases by 68.0%, 70.8%, and 98.3% for UniXcoder, CodeT5, and CoCoSoDa, respectively. Surprisingly, the simple BM25 method achieves the best performance among all RAG approaches, outperforming the second-best method CoCoSoDa by 9.3% in EM. The results suggest that in C++ code retrieval, lexical matching might be more effective than semantic matching through neural embeddings. This could be attributed to C++'s strict syntax and the prevalence of domain-specific identifiers, where exact keyword matches might be more indicative of relevance than contextual similarities captured by neural models. Moreover, we observe that RAG with CodeBERT embeddings performs poorly, showing similar results with random retrieval and even leading to slight performance degradation. This finding underscores the critical dependence of similarity-based RAG on retrieval accuracy - when the embedding model fails to effectively capture code similarities (as in the case of CodeBERT), it may adversely affect the code completion performance.

For dependency-based retrieval RAG_{Dependency}, we observe modest improvements of 8.9% in EM compared to the base model. Such gains are substantially smaller than those achieved by similarity-based retrieval, suggesting that retrieving similar implementations is more beneficial than retrieving function dependencies for line-level code completion tasks.

Fine-tuning also demonstrates notable improvements in model performance. Specifically, after fine-tuning, the models show average improvements of 78.3%, 21.3%, and 44.6% in exact match, edit similarity, and BLEU score respectively. However, we observe that the overall performance of fine-tuned models, while substantial, does not match the upper bound achieved by RAG_{sim}-BM25 and

RAG_{sim}-CoCoSoDa. Nevertheless, fine-tuning does outperform RAG_{sim} variants using other embedding models, positioning it as a competitive adaptation strategy. This comparison suggests that while fine-tuning offers consistent improvements, the quality of retrieved examples in RAG approaches can potentially lead to superior performance when optimal embedding models are employed.

Finding 1: Our evaluation of LCMs on industrial code completion reveals that LCMs demonstrate limited effectiveness on industrial code completion, indicating a substantial domain gap. In addition, both RAG and fine-tuning substantially improve performance in all evaluation metrics. RAG approaches achieve a higher performance ceiling compared to fine-tuning, with BM25 emerging as the most effective method compared to neural embedding models.

4.2 RQ2: Combining RAG and Fine-Tuning

To investigate whether fine-tuning can be effectively combined with RAG, we conduct experiments using different retrieval methods on fine-tuned and base models. The results are shown in Table 4. Note that the following experiments we focus on RAG_{sim} due to their superior performance.

Our experiments demonstrate that combining fine-tuning with RAG consistently yields better performance across all retrieval methods. For instance, using BM25 as the retrieval method, fine-tuned models with RAG achieve improvements of 3.67, 2.48, and 4.21 absolute percentage points in EM, ES, and BLEU scores respectively, compared to base models with RAG. This performance boost is even more pronounced with other retrieval methods. Overall, the average improvements on the three metrics reach 7.79, 5.27, and 7.91, respectively.

These consistent improvements across different retrieval methods suggest that fine-tuning and RAG complement each other effectively. The domain adaptation achieved through fine-tuning appears to enhance the model's ability to better utilize retrieved contexts, resulting in superior performance compared to using either approach independently.

Finding 2: Combining fine-tuning and RAG creates synergistic effects in improving code completion performance. When used together, these approaches consistently outperform their individual applications, with average improvements of 7.79%, 5.27%, and 7.91% in exact match, edit similarity, and BLEU scores respectively.

4.3 RQ3: Efficiency of RAG and Fine-Tuning

According to previous work [60], code completion is a time-sensitive task, leading efficiency to a crucial factor when adopting RAG and fine-tuning in this task. Therefore, in this section, we quantitatively investigate the efficiency impact of RAG and fine-tuning approaches on code completion. Our analysis focuses on both the **preparation stage** and **runtime stage** to examine their time and resource consumption.

4.3.1 Preparation Stage. For the **preparation stage**, the time and resource consumption are presented in Table 5. Fine-tuning primarily consumes GPU resources, and due to the large volume of training

Table 3: Results of different methods in the code completion task, where QC and DSC denote Qwen2.5-Coder and DeepSeek-Coder, respectively. The number of retrieved code snippets of $RAG_{sim}(K)$ is 5.

Method	QC-0.5B	QC-1.5B	QC-3B	QC-7B	DSC-1.3B	DSC-6.7B	Avg	Improve
Exact Match								
Base	21.09	24.48	25.89	28.04	23.16	26.07	24.79	-
Random	20.40	23.72	25.09	27.31	22.52	25.89	24.16	↓ 2.5%
RAG_{sim} -BM25	52.52	55.81	57.13	58.08	48.90	50.13	53.76	↑ 116.9%
RAG_{sim} -CodeBERT	20.48	23.79	25.29	27.20	22.36	25.86	24.16	↓ 2.5%
RAG_{sim} -UniXcoder	38.39	41.47	43.16	46.99	38.98	40.46	41.64	↑ 68.0%
RAG_{sim} -CodeT5	39.05	42.38	43.91	45.42	40.37	42.93	42.34	↑ 70.8%
RAG_{sim} -CoCoSoDa	45.95	49.35	50.90	52.28	47.03	49.49	49.17	↑ 98.3%
$RAG_{Dependency}$	23.00	27.05	28.57	30.11	25.16	28.09	26.99	↑ 8.9%
Fine-tuning	40.82	44.63	45.89	50.48	39.64	43.75	44.20	↑ 78.3%
Edit Similarity								
Base	58.12	60.78	62.10	63.64	57.98	60.60	60.54	-
Random	57.45	60.06	61.34	62.97	57.57	60.12	59.92	↓ 1.0%
RAG_{sim} -BM25	75.44	77.26	78.37	78.94	67.86	68.64	74.42	↑ 22.9%
RAG_{sim} -CodeBERT	57.48	60.21	61.63	62.85	57.29	60.14	59.93	↓ 1.0%
RAG_{sim} -UniXcoder	67.48	70.08	71.57	74.95	68.02	70.48	70.43	↑ 16.3%
RAG_{sim} -CodeT5	68.18	70.49	71.81	72.87	68.49	70.56	70.40	↑ 16.3%
RAG_{sim} -CoCoSoDa	72.17	74.56	75.67	76.81	72.49	74.45	74.36	↑ 22.8%
$RAG_{Dependency}$	59.98	61.62	63.79	65.08	59.24	62.10	61.97	↑ 2.4%
Fine-tuning	71.52	74.12	74.82	77.28	70.10	72.87	73.45	↑ 21.3%
BLEU Score								
Base	38.15	41.47	43.02	44.73	38.44	41.63	41.24	-
Random	37.82	41.03	42.56	44.37	38.74	41.20	40.95	↓ 0.7%
RAG_{sim} -BM25	66.22	68.35	69.64	70.34	58.87	59.99	65.57	↑ 59.0%
RAG_{sim} -CodeBERT	37.93	41.05	42.79	44.29	38.33	41.05	40.91	↓ 0.8%
RAG_{sim} -UniXcoder	54.55	55.77	57.30	62.22	52.23	55.19	56.21	↑ 36.3%
RAG_{sim} -CodeT5	53.78	56.75	58.08	59.41	53.87	56.62	56.42	↑ 36.8%
RAG_{sim} -CoCoSoDa	59.71	62.61	63.77	64.99	59.64	62.05	62.13	↑ 50.7%
$RAG_{Dependency}$	41.27	44.55	46.26	47.86	41.10	44.73	44.29	↑ 7.4%
Fine-tuning	57.02	60.50	61.69	65.17	54.82	58.51	59.62	↑ 44.6%

Table 4: Results of combining RAG with fine-tuning, where the results are the average of six experimented LCMs.

Methods	EM	ES	BLEU
Base+ RAG_{sim} -BM25	53.76	74.42	65.57
FT+ RAG_{sim} -BM25	57.43	76.90	69.78
Base+ RAG_{sim} -UniXcoder	41.64	70.43	56.21
FT+ RAG_{sim} -UniXcoder	52.08	77.05	66.45
Base+ RAG_{sim} -CodeT5	42.34	70.40	56.42
FT+ RAG_{sim} -CodeT5	52.11	77.18	66.05
Base+ RAG_{sim} -CoCoSoDa	49.17	74.36	62.13
FT+ RAG_{sim} -CoCoSoDa	56.45	79.55	69.68

data and model size, the training process is computationally intensive as shown in Table 5. On our 8×A100 80GB GPU server, training Qwen2.5-Coder requires 3.5 to 23.2 hours, varying by model size from 0.5B to 7B parameters, with GPU memory usage ranging from

18.3 GB to 74.8 GB. Similarly, training DeepSeek-Coder 1.3B and 6.7B consumes 6.8 and 41.4 hours respectively, utilizing 17.6 GB and 67.2 GB of GPU memory. These substantial resource requirements and A100 hours represent a remarkable bottleneck for the fine-tuning approach.

For RAG approaches, the preparation stage primarily consists of two phases: embedding computation and index construction. As shown in Table 6, the time consumption varies greatly across different methods. BM25, being a lexical-based method, requires no embedding computation and completes index construction in merely 2 minutes. In contrast, neural embedding-based methods incur substantial computational overhead. The embedding phase, executed on a single A100 GPU, takes between 34 minutes (CodeT5) and 74 minutes (UniXcoder). The subsequent indexing phase for these neural methods requires an additional 19-37 minutes, with UniXcoder taking the longest at 37 minutes. In total, while BM25 completes preparation in 2 minutes, neural embedding-based approaches require 62-111 minutes. The variation in computation

Table 5: Preparation (training) time and resource consumption of fine-tuning different LCMs.

	QC-0.5B	QC-1.5B	QC-3B	QC-7B	DSC-1.3B	DSC-6.7B
Time	3.5h	4.6h	10.1h	23.2h	6.8h	41.4h
Memory	18.3G	31.5G	44.5G	74.8G	17.6G	67.2G

time among neural models primarily stems from their architectural differences and embedding dimensions.

4.3.2 Runtime Stage. For the runtime stage, the efficiency impact of RAG can be attributed to two main factors: (1) the embedding and retrieval phase of code completion queries (using embedding models), and (2) the increased query length due to the concatenation of Top-k retrieved code snippets, which subsequently affects model inference speed.

For the first factor, we measure the time consumption of both embedding and retrieval phases across our five similarity-based RAG methods, with results shown in Table 6. We find that during runtime, BM25 demonstrates superior efficiency with only 10.8ms for retrieval operations, obviously less than neural embedding-based methods. Among neural approaches, CodeBERT and UniXcoder show moderate latency (49.2ms and 46.8ms), while CodeT5 and CoCoSoDa require substantially more time (116ms and 114ms).


For the second factor, we conduct an in-depth analysis using BM25 as a representative case to investigate how RAG affects query length and model inference efficiency. As shown in Figure 1, we compare the average input length and model throughput across various scenarios: from base queries to queries augmented with top-1 through top-5 similar code snippets. The base queries have an average length of 246 tokens, but this length increases linearly with each additional retrieved code snippet, reaching approximately 2,500 tokens when incorporating top-5 similar codes. This increase in input length significantly impacts the model’s throughput. Taking BM25 as an example, the average processing speed drops substantially from 4,031 tokens/second for base queries to around 874 tokens/second when including top-5 retrieved snippets.

In contrast, fine-tuning introduces no additional efficiency overhead during the runtime stage. After the resource-intensive fine-tuning process, where model parameters are adapted to our domain-specific dataset, the model can achieve substantial performance improvements using only base queries during inference. This makes fine-tuning more efficient at runtime compared to RAG, despite its higher upfront computational costs during the preparation stage.

Finding 3: RAG and fine-tuning exhibit distinct efficiency trade-offs in code completion. Fine-tuning requires substantial computational resources (e.g., 41.4 hours and 67.2 GB GPU memory for DeepSeek-Coder 6.7B) but introduces no runtime overhead. In contrast, RAG has minimal preparation costs but incurs runtime overhead through context retrieval and increased input lengths, reducing model throughput by up to 78% (from 4,031 to 874 tokens/second) when using top-5 retrieved snippets.

4.4 RQ4: Parameter Analysis

4.4.1 Impact of Retrieved Context Size. We conduct a detailed analysis of how base models and fine-tuned models perform when


Figure 1: Average token numbers and throughput under different retrieved code snippets across six LCMs.
Table 6: Preparation (embedding and indexing) and retrieval time consumption of RAG methods.

Stage	BM25	CodeBERT	UniXcoder	CodeT5	CoCoSoDa
Embedding	-	53min	74min	34min	64min
Indexing	2min	19min	37min	28min	35min
Runtime Retrieval	10.8ms	49.2ms	46.8ms	116ms	114ms

combined with RAG, specifically examining performance changes as the number of retrieved contexts (top-k) increases. We exclude CodeBERT from this analysis due to its poor performance in selecting relevant contexts. As shown in Figure 2, both base models and fine-tuned models demonstrate performance improvements as the number of retrieved contexts (K) increases. However, the improvement patterns differ significantly between them. While fine-tuned models combined with RAG achieve higher absolute performance, the marginal benefit (slope of improvement) from additional retrieved contexts is less pronounced compared to base models with RAG.

4.4.2 Impact of the Scale of CodeBase. Both RAG and fine-tuning require a codebase during their adaptation process, i.e., either as a retrieval database or training dataset. In this research question, we investigate how sensitive these approaches are to the scale of the codebase. We randomly sample different sizes of C++ files (30k, 60k, 90k, and 120k) from our codebase to create varying-sized retrieval databases and training sets. Due to the computational intensity of training and evaluation, we conduct experiments using Qwen models of smaller sizes (0.5B, 1.5B, and 3B parameters). For RAG, we focus on BM25 and CoCoSoDa as retrieval methods, as they demonstrate superior performance in our previous experiments. The results of these experiments are presented in Figure 3.

The experimental results reveal that both fine-tuning and RAG exhibit performance improvements as the codebase size increases, with the most substantial gains observed when transitioning from the base model to the 30K files included. Specifically, across the three model sizes, fine-tuning and RAG-BM25 demonstrate average relative improvements of 64.6% and 78.3% in EM, respectively. As the codebase expands further, the performance gains gradually diminish, but with notably different patterns between fine-tuning and RAG. When scaling from 90K to 120K files, fine-tuning shows diminishing returns, with average improvements of merely 0.35% and 0.16% percentage points in exact match and edit similarity across the three models. In contrast, RAG methods maintain more substantial improvements: RAG-BM25 and RAG-CoCoSoDa achieve




Figure 2: Exact match of the base and fine-tuned model under different retrieved contexts.




Figure 3: Exact match and edit similarity of fine-tuning and RAG under different sizes of codebase.

average EM gains of 2.26% and 2.13% across all model sizes. This phenomenon suggests that fine-tuning approaches a performance plateau in our dataset when the training data exceeds approximately 300M tokens (90K files). However, RAG methods continue to derive meaningful benefits from larger-scale codebases, demonstrating better scalability in leveraging additional code resources.

Finding 4: Both RAG and fine-tuning demonstrate sensitivity to codebase scale but with distinct scaling patterns. While both approaches show substantial initial gains, fine-tuning exhibits diminishing returns beyond 90K files with minimal improvements. In contrast, RAG methods maintain meaningful performance gains even at larger scales (90K to 120K files), suggesting superior scalability in leveraging expanded codebases for code completion.

5 Discussion

5.1 Impact of Fine-tuning on Model's Generalization Ability

In this section, we evaluate the catastrophic forgetting effect on generalization ability induced by fine-tuning, specifically examining how the model's performance on other tasks and benchmarks is affected after fine-tuning on our code repository. We select three representative tasks: code generation, code reasoning, and code translation. For code generation, following the work [42, 68], we evaluate the HumanEval and MBPP datasets using pass@1 as the

evaluation metric. For code reasoning, we employ the CRUX-Eval dataset [14], where the model is required to infer either the output given code and input (CRUX-O), or the input given code and output (CRUX-I), also measured by pass@1. For code translation, we utilize the CodeXGLUE benchmark [39], measuring performance with the CodeBLEU metric [51]. The experimental results are presented in Table 7, which demonstrates the model's performance on these general tasks before and after fine-tuning on our domain-specific dataset.

Based on our experimental results, we observe varying degrees of performance degradation across different tasks, indicating a nuanced pattern of catastrophic forgetting [13, 44]. Specifically, in code generation tasks, we observe average Pass@1 decreases of 24.1% and 11.7% on HumanEval and MBPP respectively. When evaluated with stricter test cases (+Plus), this performance gap widens to 24.0% and 16.7%.

Regarding code reasoning and translation capabilities, we observe a much milder impact from fine-tuning compared to code generation tasks. For code reasoning, fine-tuned models show average decreases of 6.5% and 5.8% on CRUX-O and CRUX-I respectively. In code translation tasks, the impact is even smaller, with only a marginal decrease of 0.7% in C#2Java and, notably, a slight improvement of 0.5% in Java2C# after fine-tuning.

This pattern suggests that fine-tuning our dataset has a minimal negative impact on the model's reasoning and translation

Table 7: Comparison of base and fine-tuned models on the code generation, code reasoning, and code translation tasks.

Models	HumanEval	+Plus	MBPP	+Plus	CRUX-O	CRUX-I	C#2Java	Java2C#
	Pass@1				CodeBLEU			
QC-0.5B	26.83	22.56	50.88	44.11	23.62	22.25	41.71	32.79
Fine-tuned	18.29	15.24	36.09	29.82	21.25	20.00	40.57	33.83
QC-1.5B	46.34	37.80	65.16	53.13	35.00	25.12	45.95	36.15
Fine-tuned	21.95	17.68	61.65	49.87	30.50	23.12	44.90	34.94
QC-3B	49.39	39.02	67.17	53.88	39.12	30.75	47.61	37.19
Fine-tuned	40.85	32.32	64.66	51.88	36.50	28.38	47.32	36.70
QC-7B	62.20	52.44	70.43	57.89	50.12	36.00	48.00	38.92
Fine-tuned	45.73	40.24	36.09	29.82	45.63	33.38	47.52	38.11
DSC-1.3B	31.10	25.61	55.89	45.86	29.88	24.75	44.80	31.22
Fine-tuned	29.27	23.78	55.39	44.36	28.25	22.62	43.93	34.57
DSC-6.7B	52.44	45.12	64.91	51.63	42.12	33.38	46.24	34.29
Fine-tuned	47.56	39.63	62.16	49.62	43.38	35.38	45.88	35.33
Avg. Base	44.72	37.04	62.41	51.08	36.64	28.71	45.72	35.09
Avg. Fine-tuned	33.95	28.15	55.14	42.56	34.25	27.15	45.02	35.60

capabilities, starkly contrasting to the more substantial degradation observed in code generation tasks. Particularly encouraging is the case of Java2C# translation, where fine-tuned models actually outperform their base counterparts on average. We also observe a positive outlier with DeepSeek-Coder 6.7B, which shows improvements in both reasoning (3.0% and 6.0% on CRUX-O and CRUX-I respectively) and translation tasks after fine-tuning. These findings suggest that fine-tuning has varying impacts across different code-related capabilities. While code generation performance shows notable degradation, the effects on code reasoning and translation abilities are much less pronounced.

5.2 Implication of Findings

In this section, we discuss the implications of our work for developers and researchers.

5.2.1 For Developers. Our main results demonstrate that RAG and fine-tuning present distinct trade-offs in effectiveness and efficiency for practical applications. Based on these findings, we offer the following implications for developers adapting LLMs in industrial settings:

(1) **Exploring fine-grained source code segmentation for RAG:** Our experiments, which used function-level retrieval, showed that adding a single relevant code snippet during inference increases input tokens by approximately 500 on average, which remarkably affects inference speed. Investigating finer-grained RAG approaches that reduce token length while maintaining retrieval effectiveness could lead to more efficient solutions. This might include exploring approaches such as snippet-level or statement-level retrieval.

(2) **Selecting RAG or fine-tuning based on available resources:** RAG and fine-tuning consume computational resources at different stages of deployment. While RAG requires more computational resources during runtime, fine-tuning demands substantial resources during the preparation stage. Given that training and inference typically utilize different GPU configurations [23, 52], developers should select their adaptation approach based on their

available resource distribution. In resource-rich environments, combining both approaches can yield optimal results.

(3) **Using RAG for ensuring LCMs' generalization ability:** When deploying models that need to handle diverse scenarios, RAG emerges as the preferred choice. Our experiments show that fine-tuning can compromise a model's generalization abilities such as in code generation and reasoning tasks. Therefore, when the deployment scenario requires maintaining both domain-specific performance and general-purpose capabilities, RAG presents a more balanced solution.

(4) **Selecting RAG in data-rich scenarios:** Our parametric studies reveal that fine-tuning's performance improvements plateau as the codebase scale increases, while RAG continues to show significant performance gains. In scenarios with abundant code data, RAG demonstrates superior scalability and can provide larger performance gains. This suggests that organizations with large proprietary codebases might find RAG particularly advantageous for their code completion systems.

5.2.2 For Researchers. In Section 4, we demonstrate that both fine-tuning and RAG are effective approaches for adapting LLMs to proprietary codebases, substantially improving code completion performance on private datasets. However, our findings reveal that each approach has distinct advantages and limitations, suggesting several promising research directions:

(1) **Exploring more advanced code embedding models:** Our experiments show that RAG's effectiveness is directly correlated with retrieval performance. Investigating more sophisticated code embedding models could lead to more relevant code retrieval and, consequently, better completion accuracy. This includes developing embedding techniques that better capture semantic relationships and domain-specific code patterns.

(2) **Exploring forgetting-aware fine-tuning techniques:** While fine-tuning effectively improves domain-specific performance, we observe degradation in general capabilities, particularly in code generation tasks. Research into more robust fine-tuning techniques that maintain the model's general capabilities while enhancing




Figure 4: Example that the QC-7B base model fails to predict the next line of code but fine-tuning and RAG succeed.

domain-specific performance represents a crucial direction. This could involve exploring methods such as elastic weight consolidation [25] or knowledge distillation [17].

(3) **Optimizing long-context inference:** We observe that inference speed is significantly affected by token length, particularly in industrial codebases which typically involve long and complex source code. Research into optimizing long-sequence handling, such as efficient attention mechanisms or context compression techniques, is crucial for improving the practical efficiency of code completion systems in industrial settings.

5.3 Case Study

In this section, we conduct case studies to illustrate how fine-tuning and RAG influence LLM’s code completion predictions in practice.

Figure 4 demonstrates the varying completion results across different approaches. For the base code completion query (shown on the right), the base QC-7B model fails to predict the correct config-related operation, generating a plausible but incorrect `Read(conf)` call. This limitation stems from the model’s lack of knowledge about the proprietary codebase’s specific APIs and conventions.

After fine-tuning, the model successfully generates the correct code snippet, specifically using the `SetFileIfUnset` method. This improvement can be attributed to the model’s acquired knowledge about `SNSLogicClientConfig` and its associated usage patterns through training on the proprietary codebase.

Similarly, the RAG-BM25 approach also generates the correct completion. The success of RAG can be traced to its retrieval mechanism, which identifies a relevant code snippet from another `Client` class (highlighted in red in the retrieved context). This retrieved code example contains similar config operations, providing a valuable reference for the current prediction task. This demonstrates how RAG effectively leverages existing codebase knowledge to guide accurate code completion.

5.4 Threat to Validity

External Validity: Our study primarily relies on an industrial codebase from our collaborating company, which may exhibit company-specific patterns and practices that differ from other software organizations. To mitigate this limitation, we deliberately collect a diverse set of code samples across different functional domains and business units within the company. Our dataset encompasses over

160,000 code files, providing a comprehensive representation of the department’s development practices and coding patterns.

Internal Validity: The efficiency measurements in our study are conducted under our available computational configurations. The relative performance characteristics of RAG and fine-tuning approaches might vary under different resource constraints. Specifically, the efficiency trade-offs could shift with varying training resources (GPU configurations), inference resources (deployment environments), or CPU resources (for retrieval operations). This suggests that our efficiency results should be interpreted within the context of our experimental setup.

6 Related Work

6.1 Retrieval-Augmented Generation

Retrieval-augmented generation (RAG) is an approach that improves generation quality by incorporating information from external knowledge sources. This approach enhances the accuracy and reliability of LLMs’ outputs and is widely used in code-related tasks [38, 65]. For example, Reacc [38] leverages both lexical copying and referring to code with similar semantics to improve code completion. Parvez et al. [47] propose REDCODER which enhances code generation and summarization by retrieving and leveraging relevant code snippets and summaries from a reference database. Similarly, Wang et al. [65] and Peng et al. [48] advanced automatic program repair by explicitly utilizing fix templates from historical bug-fix pairs.

6.2 Large Code Models

Recently, Large Code Models have emerged as powerful tools for diverse software engineering tasks [21, 22, 33, 34, 43, 61, 62, 68, 71, 74, 76]. Various influential models have demonstrated significant capabilities in code generation and understanding. For example, Meta’s Code Llama [54] is a foundation model that builds upon the original LLaMA architecture and extends the context window to 16K tokens and serves as a versatile foundation model for programming tasks. DeepSeek-Coder[16] is a series of large code models that have an architecture identical to CodeLlama. DeepSeek-Coder is trained from 2T tokens from scratch, achieving state-of-the-art performance in a variety of code intelligence tasks. Qwen2.5 Coder [18] is a family of LCMs based on models of the Qwen2.5 series. Qwen2.5 undergoes four-stage training with a total number of 5.5T tokens, presenting state-of-the-art performance in both code generation and reasoning.

6.3 Code Completion

Code completion predicts subsequent code tokens or statements to assist programmers during software development [29–32]. With the development of deep learning, various DL-based code completion methods are applied to code completion and achieve state-of-the-art performance. For instance, Li et al. [28] proposed a point-mixture network to relieve the out-of-vocabulary problem in sequence-to-sequence code completion. Izadi et al. [20] proposed CodeFill, which integrates both type and semantic information to improve the precision of completion. Recent research has increasingly focused on leveraging repository-level context. For example, Repocoder [74] utilizes useful information scattered in different files by iterative

completion and retrieval. Liu et al. [35] proposed RepoBench to evaluate code completion systems under complex and multi-file programming scenarios. RLcoder [66] employs RAG and trains the retriever to learn to retrieve useful content from the repository.

7 Conclusion

In this paper, we conduct extensive experiments to evaluate the effectiveness of RAG and fine-tuning in LCM-based code completion in the industry. We reveal the effectiveness and efficiency trade-offs when adopting the two methods: RAG achieves a higher performance ceiling while introducing non-trivial runtime completion overhead. Our findings provide guidance for both researchers and developers for further research directions and LCM adoption.

Acknowledgment

The work described in this paper was supported by the Research Grants Council of the Hong Kong Special Administrative Region, China (No. SRFS2425-4S03 of the Senior Research Fellow Scheme). This work is also supported by the CCF-Huawei Populus Grove Fund.

References

- [1] Toufique Ahmed, Christian Bird, Premkumar Devanbu, and Saikat Chakraborty. 2024. Studying LLM Performance on Closed-and Open-source Data. *arXiv preprint arXiv:2402.15100* (2024).
- [2] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6–12, 2020, virtual*.
- [3] Tianqi Chen, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. 2016. Training deep nets with sublinear memory cost. *arXiv preprint arXiv:1604.06174* (2016).
- [4] ClangD. 2025. ClangD. <https://clangd.lvm.org/>.
- [5] copilot. 2025. copilot. <https://copilot.microsoft.com/>.
- [6] cursor. 2025. cursor. <https://www.cursor.com/>.
- [7] Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. 2022. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems* 35 (2022), 16344–16359.
- [8] Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Daxin Jiang, et al. 2020. CodeBERT: A Pre-Trained Model for Programming and Natural Languages. In *Findings of the Association for Computational Linguistics: EMNLP 2020*. 1536–1547.
- [9] Shuzheng Gao, Cuiyun Gao, Wenchao Gu, and Michael Lyu. 2024. Search-based llms for code optimization. In *2025 IEEE/ACM 47th International Conference on Software Engineering (ICSE)*. IEEE Computer Society, 254–266.
- [10] Shuzheng Gao, Cuiyun Gao, Yulan He, Jichuan Zeng, Lunyiu Nie, Xin Xia, and Michael R. Lyu. 2023. Code Structure-Guided Transformer for Source Code Summarization. *ACM Trans. Softw. Eng. Methodol.* 32, 1 (2023), 23:1–23:32.
- [11] Shuzheng Gao, Wenxin Mao, Cuiyun Gao, Li Li, Xing Hu, Xin Xia, and Michael R. Lyu. 2024. Learning in the Wild: Towards Leveraging Unlabeled Data for Effectively Tuning Pre-trained Code Models. In *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering, ICSE 2024, Lisbon, Portugal, April 14–20, 2024*. ACM, 80:1–80:13.
- [12] Shuzheng Gao, Xin-Cheng Wen, Cuiyun Gao, Wenxuan Wang, Hongyu Zhang, and Michael R. Lyu. 2023. What makes good in-context demonstrations for code intelligence tasks with llms?. In *2023 38th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 761–773.
- [13] Shuzheng Gao, Hongyu Zhang, Cuiyun Gao, and Chaozheng Wang. 2023. Keeping Pace with Ever-Increasing Data: Towards Continual Learning of Code Intelligence Models. In *45th IEEE/ACM International Conference on Software Engineering, ICSE 2023, Melbourne, Australia, May 14–20, 2023*. IEEE, 30–42.
- [14] Alex Gu, Baptiste Rozière, Hugh James Leather, Armando Solar-Lezama, Gabriel Synnaeve, and Sida Wang. 2024. CRUXEval: A Benchmark for Code Reasoning, Understanding and Execution. In *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21–27, 2024*. OpenReview.net.
- [15] Daya Guo, Shuai Lu, Nan Duan, Yanlin Wang, Ming Zhou, and Jian Yin. 2022. UniXcoder: Unified Cross-Modal Pre-training for Code Representation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 7212–7225.
- [16] Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y Wu, YK Li, et al. 2024. DeepSeek-Coder: When the Large Language Model Meets Programming—The Rise of Code Intelligence. *arXiv preprint arXiv:2401.14196* (2024).
- [17] Geoffrey E. Hinton, Oriol Vinyals, and Jeffrey Dean. 2015. Distilling the Knowledge in a Neural Network. *CoRR abs/1503.02531* (2015).
- [18] Binyuan Hui, Jian Yang, Zeyu Cui, Jiayi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, et al. 2024. Qwen2. 5-coder technical report. *arXiv preprint arXiv:2409.12186* (2024).
- [19] Loshchilov Ilya and Hutter Frank. 2018. Decoupled Weight Decay Regularization. *International Conference on Learning Representations, ICLR* (2018).
- [20] Maliheh Izadi, Roberta Gismonti, and Georgios Gousios. 2022. CodeFill: Multi-token Code Completion by Jointly learning from Structure and Naming Sequences. In *44th IEEE/ACM 44th International Conference on Software Engineering, ICSE 2022, Pittsburgh, PA, USA, May 25–27, 2022*. ACM, 401–412.
- [21] Zhenlan Ji, Pingchuan Ma, Zongjie Li, and Shuai Wang. 2023. Benchmarking and explaining large language model-based code generation: A causality-centric approach. *arXiv preprint arXiv:2310.06680* (2023).
- [22] Zhenlan Ji, Daoyuan Wu, Pingchuan Ma, Zongjie Li, and Shuai Wang. 2024. Testing and Understanding Erroneous Planning in LLM Agents through Synthesized User Inputs. *arXiv preprint arXiv:2404.17833* (2024).
- [23] Ziheng Jiang, Tianqi Chen, and Mu Li. 2018. Efficient Deep Learning Inference on Edge Devices. *MLsys* (2018).
- [24] Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2019. Generalization through memorization: Nearest neighbor language models. *arXiv preprint arXiv:1911.00172* (2019).
- [25] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. 2017. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences* 114, 13 (2017), 3521–3526.
- [26] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles*. 611–626.
- [27] Guochang Li, Chen Zhi, Jialiang Chen, Junxiao Han, and Shuiguang Deng. 2024. Exploring Parameter-Efficient Fine-Tuning of Large Language Model on Automated Program Repair. In *Proceedings of the 39th IEEE/ACM International Conference on Automated Software Engineering, ASE 2024, Sacramento, CA, USA, October 27 - November 1, 2024*. ACM, 719–731.
- [28] Jian Li, Yue Wang, Michael R. Lyu, and Irwin King. 2018. Code Completion with Neural Attention and Pointer Networks. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence, IJCAI 2018, July 13–19, 2018, Stockholm, Sweden*. ijcai.org, 4159–4165.
- [29] Zongjie Li, Pingchuan Ma, Huaijin Wang, Shuai Wang, Qiyi Tang, Sen Nie, and Shi Wu. 2022. Unleashing the Power of Compiler Intermediate Representation to Enhance Neural Program Embeddings. In *44th IEEE/ACM 44th International Conference on Software Engineering, ICSE 2022, Pittsburgh, PA, USA, May 25–27, 2022*. ACM.
- [30] Zongjie Li, Chaozheng Wang, Zhibo Liu, Haoxuan Wang, Dong Chen, Shuai Wang, and Cuiyun Gao. 2023. CCTEST: Testing and Repairing Code Completion Systems. In *45th IEEE/ACM International Conference on Software Engineering, ICSE 2023, Melbourne, Australia, May 14–20, 2023*. IEEE, 1238–1250.
- [31] Zongjie Li, Chaozheng Wang, Zhibo Liu, Haoxuan Wang, Dong Chen, Shuai Wang, and Cuiyun Gao. 2023. CCTEST: Testing and Repairing Code Completion Systems. In *45th IEEE/ACM International Conference on Software Engineering, ICSE 2023, Melbourne, Australia, May 14–20, 2023*. IEEE, 1238–1250.
- [32] Zongjie Li, Chaozheng Wang, Pingchuan Ma, Chaowei Liu, Shuai Wang, Daoyuan Wu, Cuiyun Gao, and Yang Liu. 2024. On Extracting Specialized Code Abilities from Large Language Models: A Feasibility Study. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering (Lisbon, Portugal) (ICSE '24)*. Association for Computing Machinery, New York, NY, USA, Article 74, 13 pages.
- [33] Zongjie Li, Chaozheng Wang, Shuai Wang, and Gao Cuiyun. 2023. Protecting Intellectual Property of Large Language Model-Based Code Generation APIs via Watermarks. In *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security, CCS 2023, Copenhagen, Denmark, November 26–30, 2023*.
- [34] Zongjie Li, Daoyuan Wu, Shuai Wang, and Zhendong Su. 2024. API-guided Dataset Synthesis to Finetune Large Code Models. *arXiv preprint arXiv:2408.08343* (2024).
- [35] Tianyang Liu, Canwen Xu, and Julian J. McAuley. 2024. RepoBench: Benchmarking Repository-Level Code Auto-Completion Systems. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024*. OpenReview.net.
- [36] Wei Liu, Ailun Yu, Daoguang Zan, Bo Shen, Wei Zhang, Haiyan Zhao, Zhi Jin, and Qianxiang Wang. 2024. GraphCoder: Enhancing Repository-Level Code Completion via Coarse-to-fine Retrieval Based on Code Context Graph. In *Proceedings of the 39th IEEE/ACM International Conference on Automated Software*

- Engineering, ASE 2024, Sacramento, CA, USA, October 27 - November 1, 2024. ACM, 570–581.
- [37] Ilya Loshchilov and Frank Hutter. 2016. SGDR: Stochastic Gradient Descent with Warm Restarts. In *International Conference on Learning Representations*.
 - [38] Shuai Lu, Nan Duan, Hoojae Han, Daya Guo, Seung-won Hwang, and Alexey Svyatkovskiy. 2022. ReACC: A Retrieval-Augmented Code Completion Framework. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, ACL 2022, Dublin, Ireland, May 22–27, 2022. Association for Computational Linguistics, 6227–6240.
 - [39] Shuai Lu, Daya Guo, Shuo Ren, Junjie Huang, Alexey Svyatkovskiy, Ambrosio Blanco, Colin Clement, Dawn Drain, Daxin Jiang, Duyu Tang, et al. [n.d.]. CodeXGLUE: A Machine Learning Benchmark Dataset for Code Understanding and Generation. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1)*.
 - [40] Shuai Lu, Daya Guo, Shuo Ren, Junjie Huang, Alexey Svyatkovskiy, Ambrosio Blanco, Colin B. Clement, Dawn Drain, Daxin Jiang, Duyu Tang, Ge Li, Lidong Zhou, Linjun Shou, Long Zhou, Michele Tufano, Ming Gong, Ming Zhou, Nan Duan, Neel Sundaresan, Shao Kun Deng, Shengyu Fu, and Shujie Liu. 2021. CodeXGLUE: A Machine Learning Benchmark Dataset for Code Understanding and Generation. (2021).
 - [41] Xing Han Lù. 2024. Bm25s: Orders of magnitude faster lexical search via eager sparse scoring. *arXiv preprint arXiv:2407.03618* (2024).
 - [42] Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, and Daxin Jiang. 2024. WizardCoder: Empowering Code Large Language Models with Evol-Instruct. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024*. OpenReview.net.
 - [43] Pingchuan Ma, Zhaoyu Wang, Zongjie Li, Zhenlan Ji, Ao Sun, Juergen Rahmel, and Shuai Wang. 2025. Reeq: Testing and Mitigating Ethically Inconsistent Suggestions of Large Language Models with Reflective Equilibrium. (2025). <https://doi.org/10.1145/3722554>
 - [44] Michael McCloskey and Neal J Cohen. 1989. Catastrophic interference in connectionist networks: The sequential learning problem. In *Psychology of learning and motivation*. Vol. 24. Elsevier, 109–165.
 - [45] Noor Nashid, Taha Shabani, Parsa Alian, and Ali Mesbah. 2024. Contextual API Completion for Unseen Repositories Using LLMs. *arXiv preprint arXiv:2405.04600* (2024).
 - [46] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a Method for Automatic Evaluation of Machine Translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics, July 6–12, 2002, Philadelphia, PA, USA*. ACL, 311–318.
 - [47] Md. Rizwan Parvez, Wasi Uddin Ahmad, Saikat Chakraborty, Baishakhi Ray, and Kai-Wei Chang. 2021. Retrieval Augmented Code Generation and Summarization. In *Findings of the Association for Computational Linguistics: EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 16–20 November, 2021*. Association for Computational Linguistics, 2719–2734.
 - [48] Yun Peng, Shuzheng Gao, Cuiyun Gao, Yintong Huo, and Michael R. Lyu. 2024. Domain Knowledge Matters: Improving Prompts with Fix Templates for Repairing Python Type Errors. In *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering, ICSE 2024, Lisbon, Portugal, April 14–20, 2024*. ACM, 4:1–4:13.
 - [49] Lutz Prechelt. 2000. An empirical comparison of c, c++, java, perl, python, rexx and tcl. *IEEE Computer* 33, 10 (2000), 23–29.
 - [50] qdrant. 2025. qdrant. <https://qdrant.tech/>.
 - [51] Shuo Ren, Daya Guo, Shuai Lu, Long Zhou, Shujie Liu, Duyu Tang, Neel Sundaresan, Ming Zhou, Ambrosio Blanco, and Shuai Ma. 2020. Codebleu: a method for automatic evaluation of code synthesis. *arXiv preprint arXiv:2009.10297* (2020).
 - [52] Albert Reuther, Peter Michaleas, Michael Jones, Vijay Gadepally, Siddharth Samsi, and Jeremy Kepner. 2022. AI and ML accelerator survey and trends. In *2022 IEEE High Performance Extreme Computing Conference (HPEC)*. IEEE, 1–10.
 - [53] Stephen Robertson, Hugo Zaragoza, et al. 2009. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends® in Information Retrieval* 3, 4 (2009), 333–389.
 - [54] Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Romain Sauvestre, Tal Remez, et al. 2023. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950* (2023).
 - [55] Ensheng Shi, Yanlin Wang, Wenchao Gu, Lun Du, Hongyu Zhang, Shi Han, Dongmei Zhang, and Hongbin Sun. 2023. Cocosoda: Effective contrastive learning for code search. In *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*. IEEE, 2198–2210.
 - [56] Ze Tang, Jidong Ge, Shangqing Liu, Tingwei Zhu, Tongtong Xu, Liguang Huang, and Bin Luo. 2023. Domain adaptive code completion via language models and decoupled domain databases. In *2023 38th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 421–433.
 - [57] tantivy. 2025. tantivy. <https://github.com/quickwit-oss/tantivy>.
 - [58] tree sitter. 2025. tree-sitter. <https://tree-sitter.github.io/tree-sitter/>.
 - [59] Chaozheng Wang, Shuzheng Gao, Cuiyun Gao, Wenxuan Wang, Chun Yong Chong, Shan Gao, and Michael R. Lyu. 2024. A Systematic Evaluation of Large Code Models in API Suggestion: When, Which, and How. In *Proceedings of the 39th IEEE/ACM International Conference on Automated Software Engineering, ASE 2024, Sacramento, CA, USA, October 27 - November 1, 2024*. ACM, 281–293.
 - [60] Chaozheng Wang, Junhao Hu, Cuiyun Gao, Yu Jin, Tao Xie, Hailiang Huang, Zhenyu Lei, and Yuetang Deng. 2023. How Practitioners Expect Code Completion?. In *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering, ESEC/FSE 2023, San Francisco, CA, USA, December 3–9, 2023*. ACM, 1294–1306.
 - [61] Chaozheng Wang, Zongjie Li, Cuiyun Gao, Wenxuan Wang, Ting Peng, Hailiang Huang, Yuetang Deng, Shuai Wang, and Michael R Lyu. 2024. Exploring Multi-Lingual Bias of Large Code Models in Code Generation. *arXiv preprint arXiv:2404.19368* (2024).
 - [62] Chaozheng Wang, Zongjie Li, Yun Peng, Shuzheng Gao, Siron Chen, Shuai Wang, Cuiyun Gao, and Michael R. Lyu. 2023. REEF: A Framework for Collecting Real-World Vulnerabilities and Fixes. In *38th IEEE/ACM International Conference on Automated Software Engineering, ASE 2023, Luxembourg, September 11–15, 2023*. IEEE, 1952–1962.
 - [63] Chaozheng Wang, Yuanhang Yang, Cuiyun Gao, Yun Peng, Hongyu Zhang, and Michael R. Lyu. [n.d.]. No more fine-tuning? an experimental evaluation of prompt tuning in code intelligence. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering, ESEC/FSE 2022, Singapore, Singapore, November 14–18, 2022*. 382–394.
 - [64] Chaozheng Wang, Yuanhang Yang, Cuiyun Gao, Yun Peng, Hongyu Zhang, and Michael R. Lyu. 2023. Prompt Tuning in Code Intelligence: An Experimental Evaluation. *IEEE Trans. Software Eng.* 49, 11 (2023), 4869–4885.
 - [65] Weishi Wang, Yue Wang, Shafiq Joty, and Steven C. H. Hoi. 2023. RAP-Gen: Retrieval-Augmented Patch Generation with CodeT5 for Automatic Program Repair. In *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering, ESEC/FSE 2023, San Francisco, CA, USA, December 3–9, 2023*. ACM, 146–158.
 - [66] Yanlin Wang, Daya Guo, Jiachi Chen, Ruikai Zhang, Yuchi Ma, and Zibin Zheng. 2024. RLCoder: Reinforcement Learning for Repository-Level Code Completion. *CoRR abs/2407.19487* (2024).
 - [67] Yue Wang, Weishi Wang, Shafiq Joty, and Steven CH Hoi. 2021. CodeT5: Identifier-aware Unified Pre-trained Encoder-Decoder Models for Code Understanding and Generation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 8696–8708.
 - [68] Yuxiang Wei, Zhe Wang, Jiawei Liu, Yifeng Ding, and Lingming Zhang. 2024. Magicoder: Empowering Code Generation with OSS-Instruct. In *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21–27, 2024*. OpenReview.net.
 - [69] Xin-Cheng Wen, Cuiyun Gao, Shuzheng Gao, Yang Xiao, and Michael R. Lyu. 2024. SCALE: Constructing Structured Natural Language Comment Trees for Software Vulnerability Detection. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis, ISSTA 2024, Vienna, Austria, September 16–20, 2024*. ACM, 235–247.
 - [70] F George Wilkie and B Hylands. 1998. Measuring complexity in C++ application software. *Software: Practice and Experience* 28, 5 (1998), 513–546.
 - [71] Wai Kin Wong, Huaijin Wang, Zongjie Li, Zhibo Liu, Shuai Wang, Qiyi Tang, Sen Nie, and Shi Wu. 2023. Refining decompiled c code with large language models. *arXiv preprint arXiv:2310.06530* (2023).
 - [72] Peng Xu, Wei Ping, Xianchao Wu, Chejian Xu, Zihan Liu, Mohammad Shoeybi, and Bryan Catanzaro. 2024. Chatqa 2: Bridging the gap to proprietary llms in long context and rag capabilities. *arXiv preprint arXiv:2407.14482* (2024).
 - [73] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. 2024. Qwen2. 5 Technical Report. *arXiv preprint arXiv:2412.15115* (2024).
 - [74] Fengji Zhang, Bei Chen, Yue Zhang, Jacky Keung, Jin Liu, Daoguang Zan, Yi Mao, Jian-Guang Lou, and Weizhu Chen. 2023. RepoCoder: Repository-Level Code Completion Through Iterative Retrieval and Generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6–10, 2023*. Association for Computational Linguistics, 2471–2484.
 - [75] Jian Zhang, Xu Wang, Hongyu Zhang, Hailong Sun, and Xudong Liu. 2020. Retrieval-based neural source code summarization. In *ICSE '20: 42nd International Conference on Software Engineering, Seoul, South Korea, 27 June - 19 July, 2020*. ACM, 1385–1397.
 - [76] Kunpeng Zhang, Zongjie Li, Daoyuan Wu, Shuai Wang, and Xin Xia. 2025. Low-Cost and Comprehensive Non-textual Input Fuzzing with LLM-Synthesized Input Generators. *arXiv preprint arXiv:2501.19282* (2025).